

**STSCI 4780/5780**  
**Bayesian computation — Beyond the basics**  
**(A selective survey)**

Tom Lored, CCAPS & SDS, Cornell University

© 2022-04-21

# Notation focusing on computational tasks

$$\begin{aligned} p(\theta|D, M) &= \frac{p(\theta|M)p(D|\theta, M)}{p(D|M)} \\ \Rightarrow p(\theta) &= \frac{\pi(\theta)\mathcal{L}(\theta)}{Z} = \frac{q(\theta)}{Z} = p(\theta), \end{aligned}$$

- $M$  = model specification (context)
- $D$  specifies observed data
- $\theta$  = model parameters, of dimension  $m$
- $p(\theta)$  = posterior pdf for  $\theta$
- $\pi(\theta)$  = prior pdf for  $\theta$
- $\mathcal{L}(\theta)$  = likelihood for  $\theta$  (likelihood function)
- $q(\theta) = \pi(\theta)\mathcal{L}(\theta)$  = “quasiposterior”
- $Z = p(D|M)$  = (marginal) likelihood for the model

# Key themes in advanced algorithms

- Combining multiple update algorithms
- Adaptation—*gently* breaking the Markov property
- Augmenting the parameter space (*increasing* dimensionality)

## Combining MH updates

No one class of proposal distributions works well for all problems  
→ consider combining multiple proposals hoping they'll have complimentary strengths (esp. in a “black box” toolkit)

Example: Combine RWM updates with *various step sizes*

Reversible update = proposal + M-H accept/reject step

Two valid ways to combine reversible updates:

- *Composition*: Follow one update by another (deterministically)
- *Mixing*: Randomly choose an update mechanism

Implementations:

- Fixed/cyclic scan (or sweep)
- Random scan
- Random sequence scan—combines composition and mixing

For theory & examples, see Geyer's 1995 and 1998 MCMC notes

# Adaptive MCMC

A proposal distribution for MH sampling typically has *tuning parameters*,  $\psi$ : we draw a candidate from  $k_\psi(y; x)$ .

- Random-walk Metropolis: Proposal width in each direction
- Independent Metropolis: Shape of proposal (location, covariance. . .)

For MH, we can't have  $\psi$  depend on the chain history—the chain wouldn't be Markov!

Simple approach: We can tune  $\psi$  using pilot runs (perhaps during burn-in), and then fix it to preserve detailed balance

*Adaptive MCMC* finds ways to adjust  $\psi$  continuously that preserves asymptotic sampling properties

Main idea: *Vanishing adaptation*

### Example: Robust adaptive Metropolis (RAM)

Motivation: Consider random-walk metropolis (RWM), but with a proposal distribution that is multivariate normal, so it can take steps along directions aligned with the posterior

This requires:

- Finding a good covariance matrix for the MVN
- Drawing a vector of correlated steps from a MVN

MVN draws: Write the covariance matrix as  $C = SS^T$ , where  $S$  is the *Cholesky factorization* of  $C$  — a lower-diagonal matrix with positive elements

Then from current position  $X_{n-1}$ , we can propose a candidate position  $Y_n$  by drawing a vector  $U_n$  of *independent* standard normal variates, and shifting and correlating them:

$$Y_n = X_{n-1} + SU_n$$

Now the challenge is choosing  $S$

## RAM algorithm

Use Metropolis updates with a correlated multivariate proposal, *altering the covariance matrix along the chain to target a desired mean acceptance rate,  $\alpha_*$* :

1. Propose  $Y_n = X_{n-1} + S_{n-1}U_n$ , where  $U_n \sim q$  is an independent random vector, and  $S_{n-1}$  is a lower-diagonal matrix with positive elements
2. With probability  $\alpha_n \equiv \min\{1, \pi(Y_n)/\pi(X_{n-1})\}$  the step is accepted, and  $X_n = Y_n$ ; otherwise the step is rejected and  $X_n = X_{n-1}$
3. Compute an updated lower-diagonal matrix  $S_n$  via

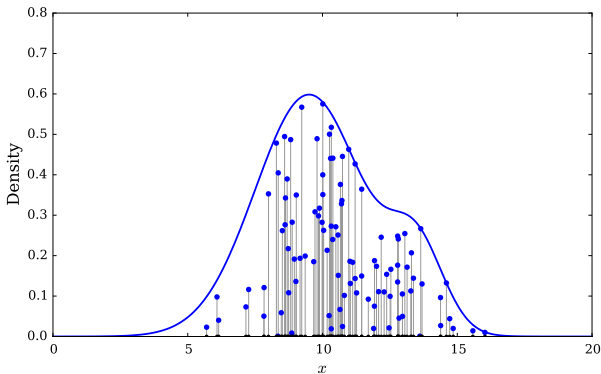
$$S_n S_n^T = S_{n-1} \left( I + \eta_n (\alpha_n - \alpha_*) \frac{U_n U_n^T}{\|U_n\|^2} \right) S_{n-1}^T \quad (1)$$

where  $I$  is an identity matrix, and  $\eta_n = n^{-2/3}$  controls the adaptivity

See: Vihola (2012): Robust adaptive Metropolis algorithm with coerced acceptance rate

# Slice sampling

We can get samples from  $p(\theta) = q(\theta)/Z$  by sampling  $(\theta, y)$  pairs uniformly in area under the  $q(\theta)$  function, and then just ignoring (marginalizing over) the  $y$  coordinates:

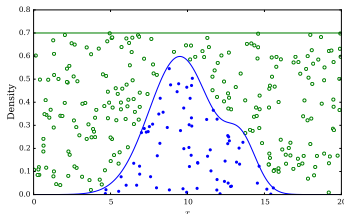




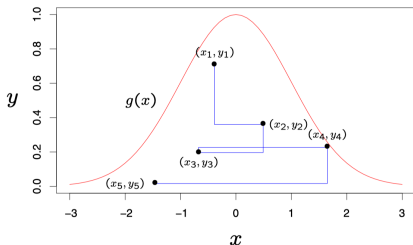
The accept-reject produces *IID samples* by sampling uniformly from an *enclosing volume*, and then *rejecting* bad samples

Slice sampling uses *Gibbs sampling* (MCMC) to sample *within* the volume, *keeping all* of the (correlated) samples)

**Accept-Reject**

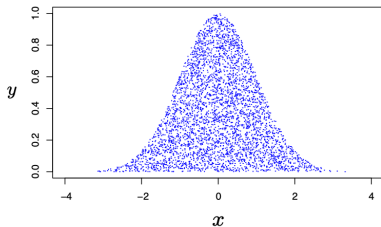


**Slice sampling**

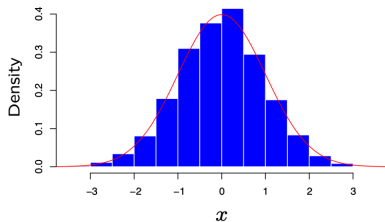


## Slice sampling for a normal dist'n:

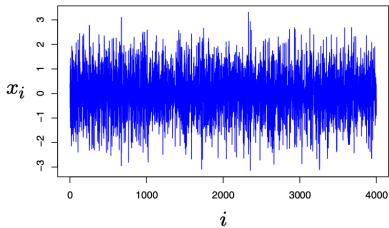
Plot of  $(x_i, y_i)$



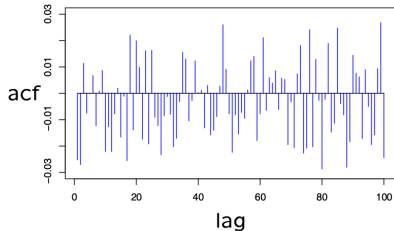
Histogram of sampled values



Trace plot



Autocorrelations



*Henderson 2000*

Slice sampling can be tricky in 1-D, and quite tricky in higher dimensions; usually implemented as one-variable-at-a-time

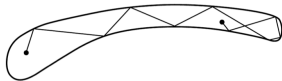
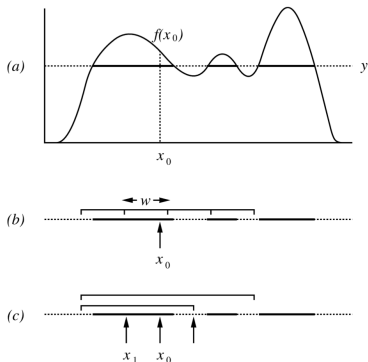


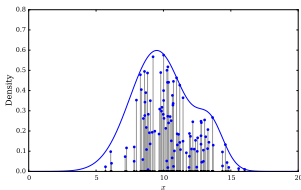
FIG. 11. Moving around a two-dimensional slice by reflection from the exact boundaries.

For details, see Slice Sampling (Neal 2000, with discussion)

## Auxiliary/augmented variables

Accept/reject and slice sampling for getting samples from a  $d$ -D density:

- Sample from a *uniform*  $(d + 1)$ -D density (with a complicated boundary):



- Report the marginal samples for the  $d$  original dimensions

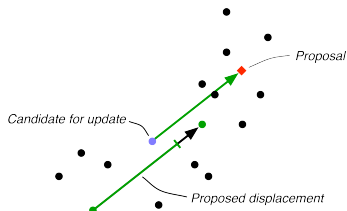
A paradoxical notion motivating some advanced MCMC methods is that making the problem “harder” (higher-dimensional) may actually make it *easier*

Hamiltonian Monte Carlo (below) and the NUTS sampler rely on auxiliary variables

# Ensemble methods: Differential Evolution MCMC

Combine evolutionary computing & MCMC (ter Braak 2006)

Follow a *ensemble/population* of states, where a randomly selected state is considered for updating via the (scaled) vector difference between two other states.



Behaves roughly like RWM, but with a proposal distribution that automatically adapts to shape & scale of posterior

Step scale: Optimal  $\gamma \approx 2.38/\sqrt{2d}$ , but occasionally switch to  $\gamma = 1$  for mode-swapping

Original DE-MCMC uses these simple moves and pop'n size  $N \sim 3d$ ; works well if given a “smart start” (initial pop'n)

Later version (ter Braak & Vrugt 2008) adds new moves and can sample effectively with just  $N = 3$  in up to a few dozen dimensions, without a smart start

A new method (not yet widely used) combines slice sampling and ensemble methods: Ensemble slice sampling

# Random Walks

Metropolis random walk (MRW) and Gibbs sampler updates execute a *random walk* through parameter space:

- Moves are local, with a characteristic scale  $l$
- Total distance traversed over time  $t \propto \sqrt{t}$

This is a relatively slow (albeit steady) rate of exploration

Multimodality  $\rightarrow$  even slower exploration; only rare large jumps can move between modes

*We need methods designed to make large moves*

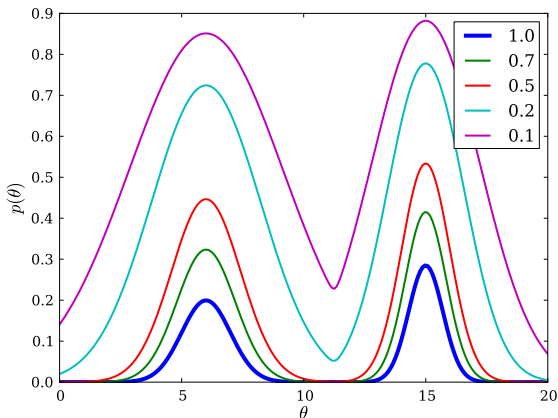
# Annealing and Parallel Tempering

*PT, aka Metropolis-coupled MCMC*

To enable large jumps, *anneal* or *temper* the posterior:

$$q_{\beta}(\theta) = [q(\theta)]^{\beta} \quad \text{or} \quad \pi(\theta)[\mathcal{L}(\theta)]^{\beta}, \quad (2)$$

$$\text{with inverse temperature/temper } \beta \in [0, 1] \quad (3)$$





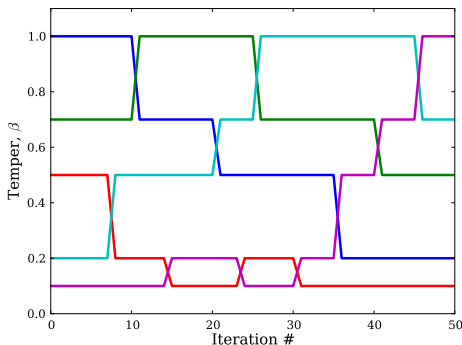
Consider a set of tempers ( “inverse temperatures” )  $\{\beta_i\}$

Think of each  $q_i = q_{\beta_i}$  as its own “model” with its own parameters, and construct a sampler for the joint distribution

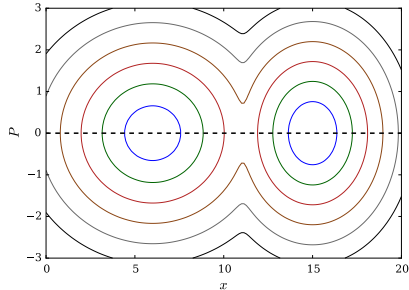
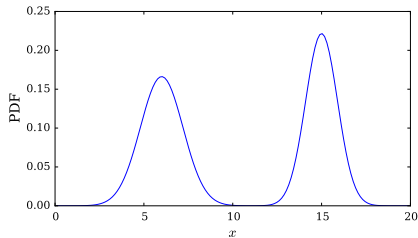
$$p(\theta_1, \dots, \theta_m) = \prod_i q_i(\theta_i)$$

Alternate within-temper proposals and swap proposals between adjacent tempers

Swaps between tempered chains



# Phase space: Doubling the dimensionality



$$p(x, P) \propto q(x) \times f(P)$$

$$p(x) = \int dP p(x, P) \propto q(x)$$

$$p(P) = \int dx p(x, P) \propto f(P)$$

- Pick  $P \sim f(P)$
- Move along a contour in phase space
- Drop  $P$ , keep  $x$

Will work if the phase space motion corresponds to sampling  $p(x, P)$

# Hamiltonian (Hybrid) Monte Carlo

Give samples “momentum” so moves tend to go in the same direction a while; use derivatives to guide the evolution → suppress random walks

Adds  $d$  additional variables,  $P$ , with a joint Gaussian dist'n:

$$\log p(\theta, P) = - \left[ U(\theta) + \frac{1}{2} P^2 \right]; \quad U(\theta) \equiv -\log q(\theta)$$

Sample  $P$  from a Gaussian, and use it to generate proposals via

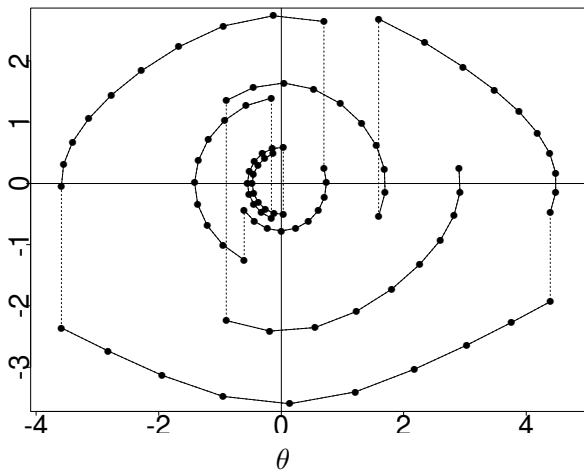
$$\dot{\theta} = P; \quad \dot{P} = -\frac{\partial H}{\partial \theta}$$

Hamiltonian dynamics → reversible, preserves volume, keeps  $p$  constant (*exact* proposals always accepted, like Gibbs sampling)

# Challenges for basic HMC

- Tuning parameters:
  - ▶ PDE integration time step size,  $\epsilon$ , and integration length,  $L$
  - ▶ Handling problems with very different scales along different dimensions ( $\rightarrow$  need different momentum scales)
- Computing the needed derivatives

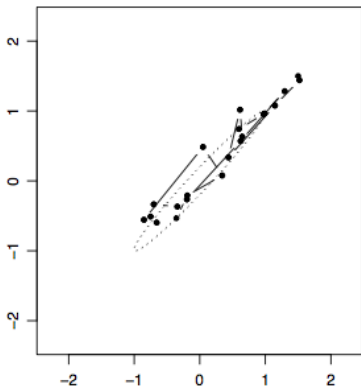
Sampling a 1-D Student- $t$  dist'n with dof= 5



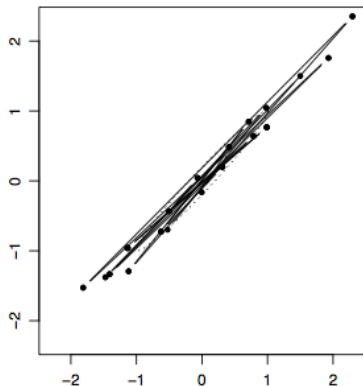
# HMC vs. random walk (2-D)

20 iterations for a correlated bivariate normal

Random-walk Metropolis



Hamiltonian Monte Carlo



Neal 2011

## Tuning integration length

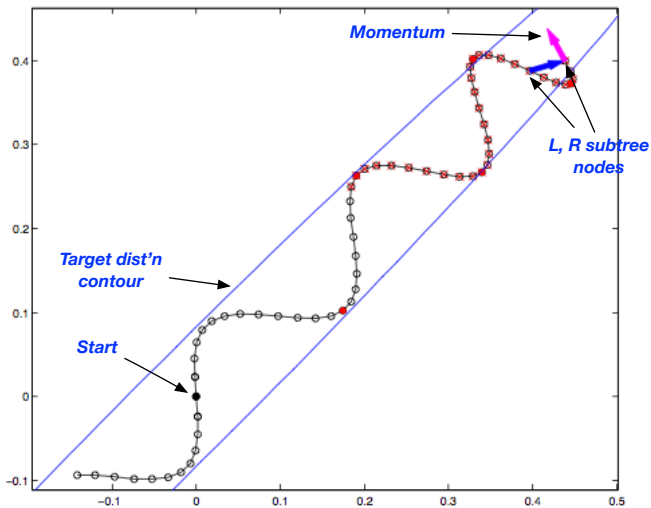
We want to move along a contour long enough to get far from the starting point, but not head back toward it

Examine rate of change of squared distance from current point,  $\theta_i$ :

$$\frac{d}{dt} \frac{(\theta - \theta_i) \cdot (\theta - \theta_i)}{2} = (\theta - \theta_i) \cdot P$$

Stop integrating when this becomes negative

## No-U-Turn Sampler (NUTS)



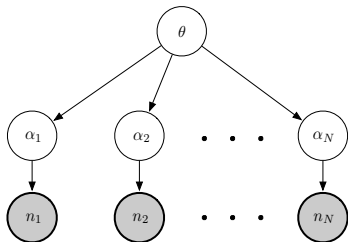
Hoffman & Gelman 2013



# Multilevel models: parameter-dependent scales

Goal: Learn a population dist'n from noisy member measurements

Qualitative



$$\begin{aligned} p(\theta, \{\alpha_i\}, \{n_i\}) &= p(\theta) \prod_i p(\alpha_i | \theta) p(n_i | \alpha_i) \\ &= \pi(\theta) \prod_i f(\alpha_i; \theta) \ell_i(\alpha_i) \end{aligned}$$

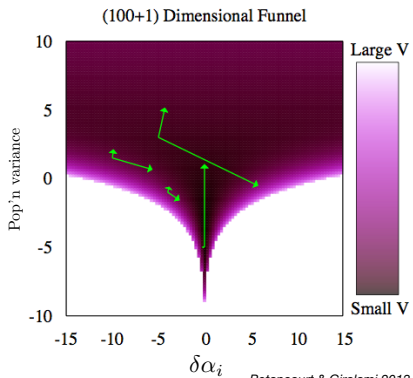
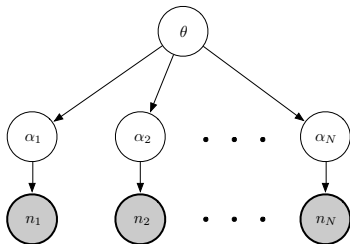
Quantitative

$$\theta = (a, b) \text{ or } (\mu, \sigma)$$

$$\pi(\theta) = \text{Flat}(\mu, \sigma)$$

$$p(\alpha_i | \theta) = \text{Beta}(\alpha_i | \theta)$$

$$p(n_i | \alpha_i) = \binom{N_i}{n_i} \alpha_i^{n_i} (1 - \alpha_i)^{N_i - n_i}$$



## Mass matrix = metric

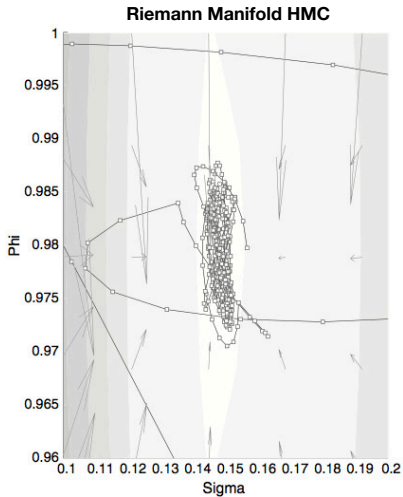
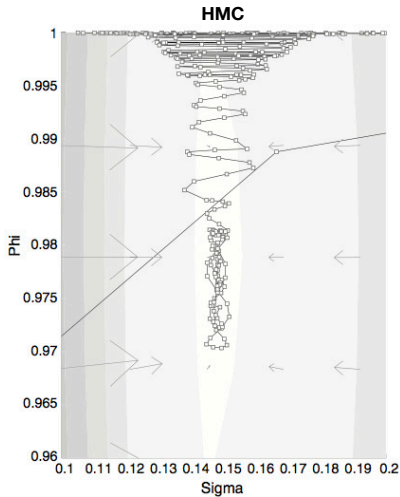
Add  $d$  additional variables,  $P$ , with a *correlated* Gaussian dist'n:

$$\log p(\theta, P) = - \left[ U(\theta) + \frac{1}{2} P \cdot M^{-1} \cdot P \right]; \quad U(\theta) \equiv -\log p(\theta)$$

$M$  introduces  $d$  more tuning parameters!

- **Euclidean manifold HMC:** Use the Hessian at the mode
- **Riemannian manifold HMC:** Use position-dependent  $M(\theta)$

# HMC vs. Riemann manifold MC



*Girolami & Calderhead 2011*

# Stan capabilities

- High-performance probabilistic model implementation
  - ▶ Stan code is compiled to a C++ library
  - ▶ Parameters transformed to unconstrained space; transformation & Jacobian handled automatically
  - ▶ Automatic differentiation (AD) used to compute derivatives of log-likelihood WRT parameters
- HMC No-U-Turn Sampler (NUTS)
  - ▶ PDE solver step size & number automatically tuned during burn-in
  - ▶ Mass matrix adaptively tuned during burn-in
- Optimization
  - ▶ BFGS and Newton's method